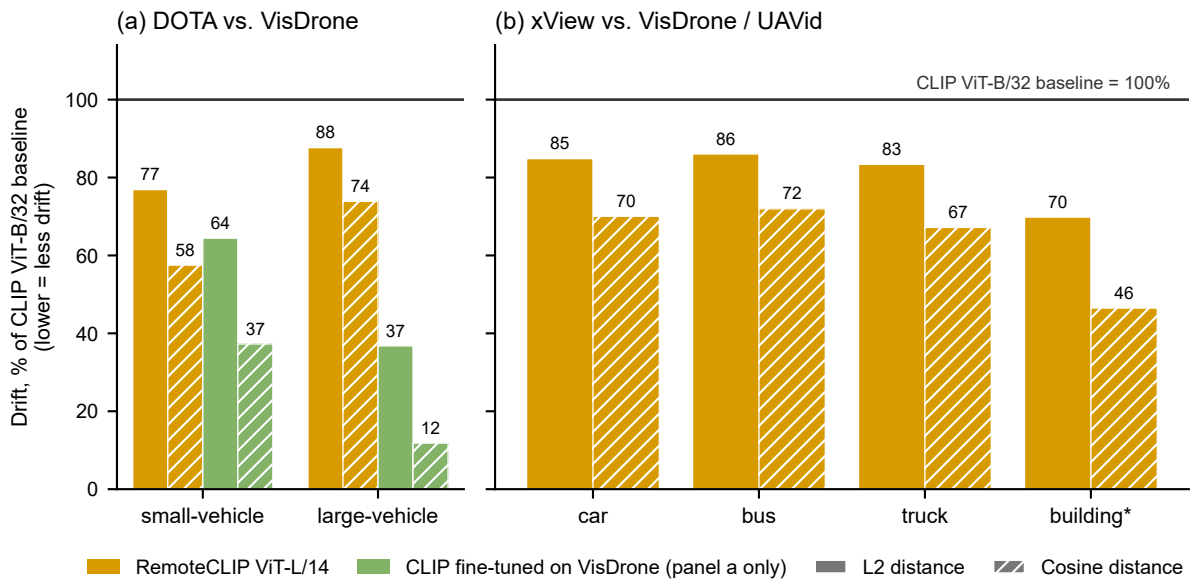

n = 300 crops per class per domain, single seed (42). All three encoders in (a) are scored on the same held-out 20% of VisDrone images; the fine-tuned model was trained on the other 80%. *building: xView vs. UAVid (VisDrone has no building class); other classes: xView vs. VisDrone.